


Unsupervised Machine Learning with Python

Section 9.0: Dimension Reduction Overview

Purpose of Dimension Reduction

- Machine learning problems may have datasets with 1000s of features (# features = # of dimensions)
- More dimensions generally means slower computation
- Dimension Reduction attempts to map feature vectors (datasets) into a lower dimensional space while retaining as much information as possible
- See [UnsupervisedML_Resources.pdf](#) for links to additional resources

Dimension Reduction



- Picture often seen in Linear Algebra courses
- Use dimension reduction techniques to find Projected Data Point and Reduced Dimensional Space

Dimension Reduction for MNIST Digits

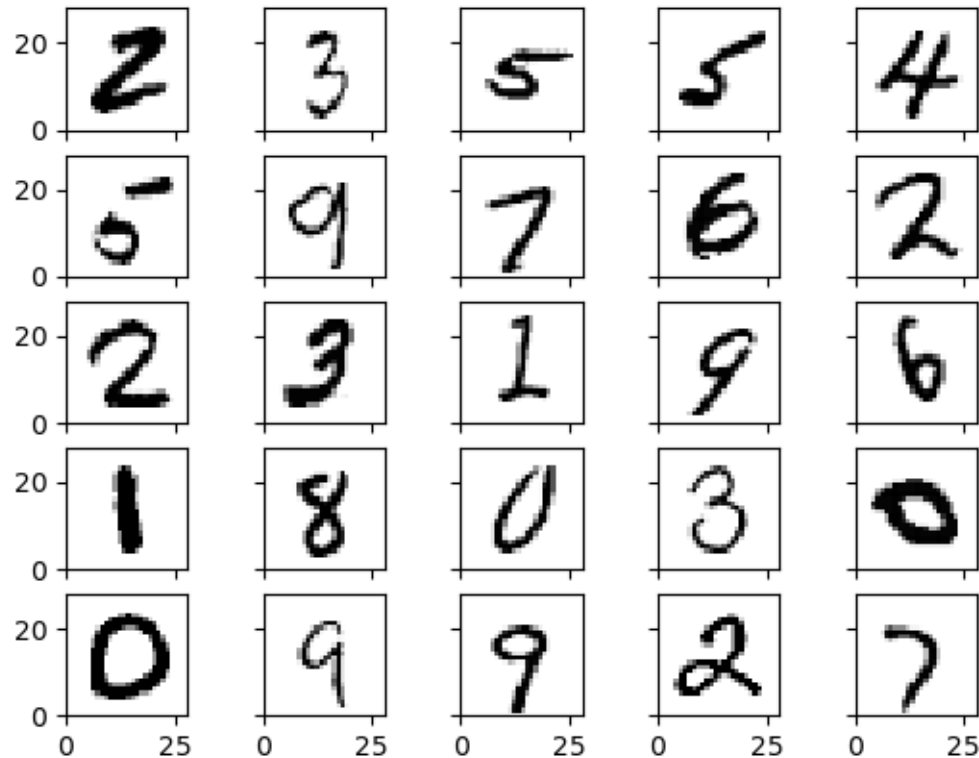
Original Dataset:

28x28 resolution = 784 pixels (dimensions)

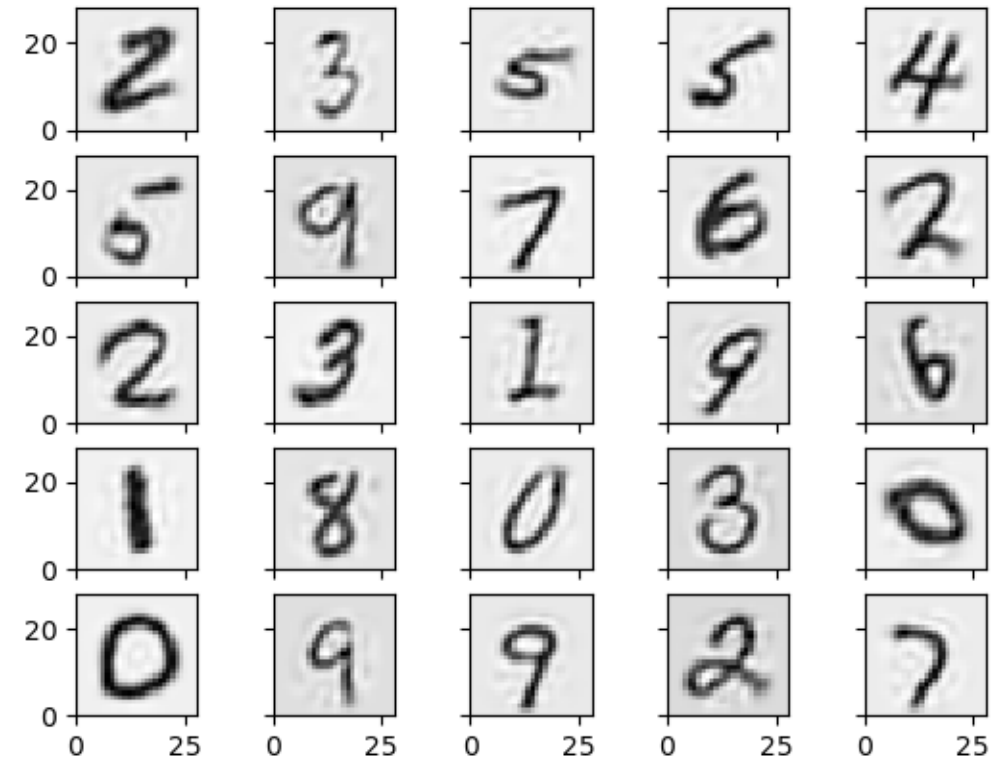
Reconstructed Dataset:

Reduce to 78 dimensions then reconstruct images

Images of Sample MNIST Digits



Images of Sample MNIST Digits



Dimension Reduction

- Principal Component Analysis (PCA):
 - “Linear” approach
 - Cool application of singular value decomposition to project a dataset onto a lower dimensional space
- Autoencoding:
 - “Nonlinear” approach
 - Uses techniques from supervised learning to learn lower dimensional representation of dataset


Unsupervised Machine Learning with Python

Section 9.1: Principal Component Analysis Algorithm

Example: 4 Data Points in 2D

- Consider 4 data points in 2 dimensions

$$X = [X_0 \quad X_1 \quad X_2 \quad X_3] = \begin{bmatrix} 1 & 1.6 & -0.5 & -1.6 \\ 0.9 & 1.65 & -0.6 & -1.5 \end{bmatrix}$$



Principal Component Analysis

Basic Idea

- Compute SVD of dataset $X = U\Sigma V^T$
- Let $U=[u_0 \ u_1 \ \dots \ u_{d-1}]$ be new coordinate system for d dimensional space
- Relevance of each basis vector $u_0 \ u_1 \ \dots \ u_{d-1}$ is determined by its corresponding singular value, which are ordered in decreasing value
- Project data onto lower dimensional space spanned by first K basis vectors u_0, u_1, \dots, u_{K-1} , choosing K to retain sufficient information

SVD of Data Set

- Compute compact SVD of X

$$X = \begin{bmatrix} 1 & 1.6 & -0.5 & -1.6 \\ 0.9 & 1.65 & -0.6 & -1.5 \end{bmatrix}$$


$$X = U\Sigma V^T = [u_0 \quad u_1] \begin{bmatrix} \sigma_0 & 0 \\ 0 & \sigma_1 \end{bmatrix} \begin{bmatrix} v_0^T \\ v_1^T \end{bmatrix}$$

$$X = \begin{bmatrix} 0.71 & -0.70 \\ 0.70 & 0.71 \end{bmatrix} \begin{bmatrix} 3.54 & 0 \\ 0 & 0.12 \end{bmatrix} \begin{bmatrix} 0.38 & 0.65 & -0.22 & -0.62 \\ -0.47 & 0.46 & -0.63 & 0.41 \end{bmatrix}$$

New Coordinate system with basis u_0 and u_1

- Consider new coordinate system using u_0 and u_1 as basis

$$u_0 = \begin{bmatrix} 0.71 \\ 0.70 \end{bmatrix} \quad u_1 = \begin{bmatrix} -0.70 \\ 0.71 \end{bmatrix}$$



Notes:

- Basis vectors of original system typically correspond to physical or measurable quantity: (word count, pixel intensity, features of a house)
- New basis vectors u_0, u_1, \dots typically don't correspond to physical or measurable quantities

Data points in u0-u1 Coordinate System

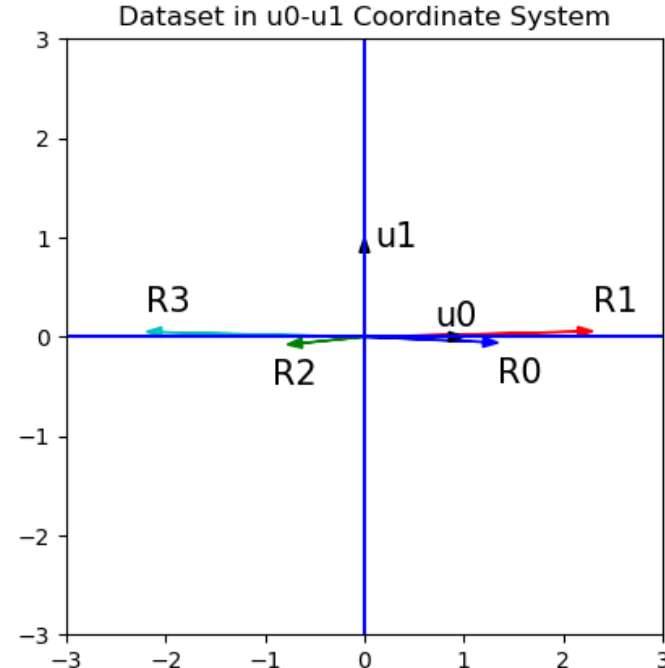
- Data points in u0 and u1 coordinate system

$$\Sigma V^T = \begin{bmatrix} 3.54 & 0 \\ 0 & 0.12 \end{bmatrix} \begin{bmatrix} 0.38 & 0.65 & -0.22 & -0.62 \\ -0.47 & 0.46 & -0.63 & 0.41 \end{bmatrix}$$

$$\Sigma V^T = [R_0 \quad R_1 \quad R_2 \quad R_3] = \begin{bmatrix} 1.34 & 2.30 & -0.78 & -2.19 \\ -0.06 & 0.06 & -0.08 & 0.05 \end{bmatrix}$$

← Units of u0 for each data point

← Units of u1 for each data point



Reduce Dimensions

- Only keep information in u0 direction (1 dimension)
- Data points in reduced number of dimensions

$$R = [\sigma_0][v_0^T] = [3.54][0.38 \quad 0.65 \quad -0.22 \quad -0.62] = [1.34 \quad 2.30 \quad -0.78 \quad -2.19]$$

Units of u0 for each data point



Reconstruct in Original Space

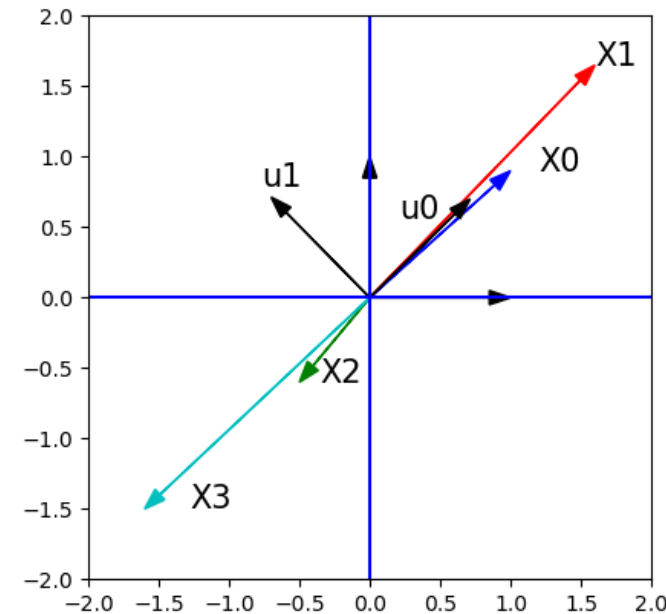
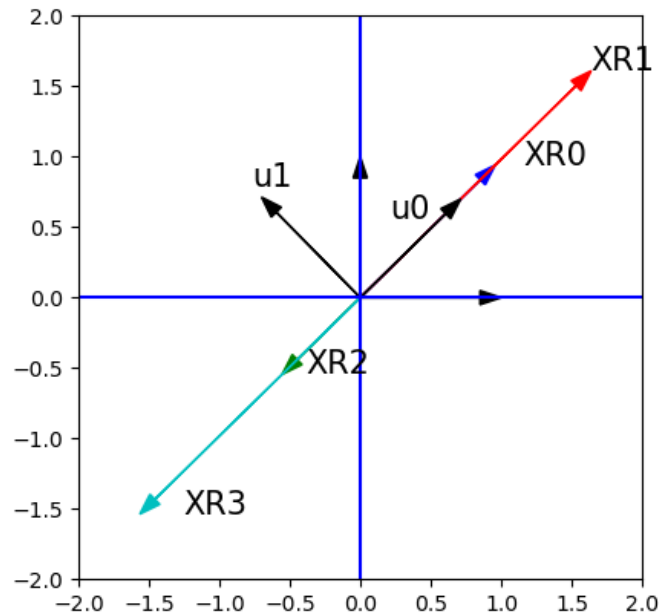
- Can reconstruct data points in original coordinate system using information in u_0 direction only

$$X_R = [u_0] R = [u_0][\sigma_0][v_0^T]$$

$$X_R = \begin{bmatrix} 0.71 \\ 0.70 \end{bmatrix} [3.54] \begin{bmatrix} 0.38 & 0.65 & -0.22 & -0.62 \end{bmatrix}$$

$$X_R = \begin{bmatrix} 0.96 & 1.64 & -0.55 & -1.56 \\ 0.94 & 1.61 & -0.54 & -1.54 \end{bmatrix}$$

$$X = \begin{bmatrix} 1 & 1.6 & -0.5 & -1.6 \\ 0.9 & 1.65 & -0.6 & -1.5 \end{bmatrix}$$



PCA Algorithm


PCA Algorithm:

- Start with data matrix X (d features \times M samples) (here $d < M$)
- (1) Subtract mean of data $\bar{X} = X - X_{mean}$ (translate data points by same amount)
- (2) Compute the compact version of SVD of $\bar{X} = X - X_{mean} = U\Sigma V^T$

$$\bar{X} (d \times M) = U (d \times d) \Sigma (d \times d) V^T (d \times M)$$

PCA Algorithm

(3) Pick $K < d$ principal components



(first K columns of U, first K diagonal entries of Σ , first K rows of V^T)

(4) Data points in K dimensional (u_0, \dots, u_{K-1}) coordinate system given by:

$$R = \begin{bmatrix} \sigma_0 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_{K-1} \end{bmatrix} \begin{bmatrix} v_0^T \\ \vdots \\ v_{K-1}^T \end{bmatrix}$$

PCA Algorithm

(5) Reconstructed \bar{X}_R in original space(keeping only K principal components)

$$\bar{X}_R = [u_0 \quad \cdots \quad u_{K-1}]R = [u_0 \quad \cdots \quad u_{K-1}] \begin{bmatrix} \sigma_0 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_{K-1} \end{bmatrix} \begin{bmatrix} v_0^T \\ \vdots \\ v_{K-1}^T \end{bmatrix}$$

(6) Add back mean to get reconstructed X (based on K principal components)

$$X_R = \bar{X}_R + X_{mean}$$

- SVD is connected to L2 distance measure theoretically, so use L2 distance measure when computing distances for consistency

Principal Component Analysis: Choosing K

- Covariance matrix for dataset:

$$Cov = \frac{1}{M} (X - X_{mean})(X - X_{mean})^T$$

- Total variance is sum of eigenvalues of Covariance matrix, which is same as sum of squares of singular values of $(X - X_{mean})$, divided by M
- Define

$$Variance(K) = \frac{1}{M} \sum_{k=0}^{K-1} \sigma_k^2$$

- Instead of specifying K directly, choose smallest number of principal components so that $Variance(K)/Variance(N)$ is at least as large as specified percentage

Example: Choosing K

- Assume $M = 4$ data points
- Assume singular values are: $\sigma_0 = 4, \sigma_1 = 3, \sigma_2 = 2, \sigma_3 = 1$
- Recall: $Variance(K) = \frac{1}{M} \sum_{k=0}^{K-1} \sigma_k^2$

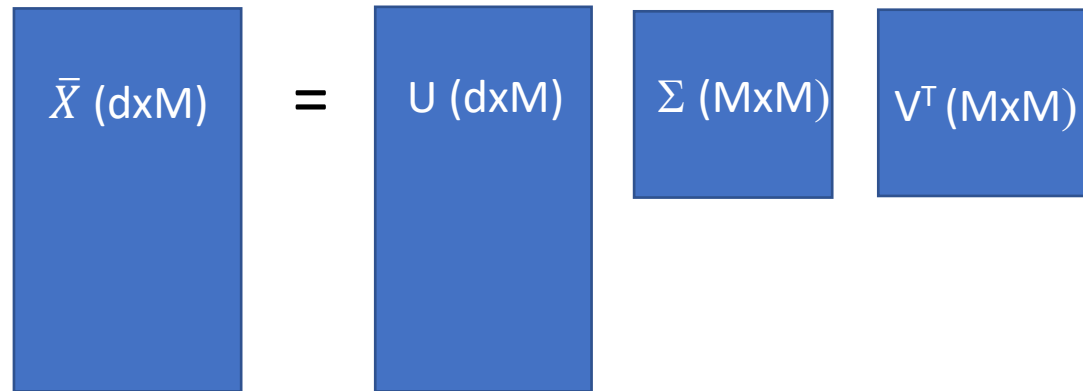
	K = 1	K = 2	K = 3	K = 4
Variance(K)	4.00	6.25	7.25	7.50
Variance(K)/Variance(M)	0.5333	0.8333	0.9666	1.0000

- In this example, only need $K=2$ principal components to capture 83% of total variance

More Dimensions than Data Points

PCA Algorithm in case number of dimensions $d >$ number of samples M

- Compute the compact version of SVD of $\bar{X} = X - X_{mean} = U\Sigma V^T$


$$\bar{X} (d \times M) = U (d \times M) \Sigma (M \times M) V^T (M \times M)$$

- Suppose we retain M principal components

$$R = \begin{bmatrix} \sigma_0 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_{M-1} \end{bmatrix} \begin{bmatrix} v_0^T \\ \vdots \\ v_{M-1}^T \end{bmatrix}$$

- R is in $M < d$ dimensions and no information is lost
- (L2) distances between data points same in original and reduced dimensional space since U has orthogonal columns of length 1.

Example

- Dataset X: 2 data points in 4 dimensions

$$X_0 = \begin{bmatrix} 1 \\ 3 \\ 5 \\ 7 \end{bmatrix} \quad X_1 = \begin{bmatrix} 2 \\ 4 \\ 6 \\ 8 \end{bmatrix} \quad X = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix}$$

- Compute L2 (Euclidean) distance between points

$$\text{dist}(X_0, X_1) = 2.646$$

Example

- SVD

$$X = U\Sigma V^T = \begin{bmatrix} -0.16 & -0.49 \\ -0.36 & -0.13 \\ -0.52 & -0.63 \\ -0.76 & 0.59 \end{bmatrix} \begin{bmatrix} 13.93 & 0 \\ 0 & 0.92 \end{bmatrix} \begin{bmatrix} -0.62 & -0.79 \\ 0.79 & -0.62 \end{bmatrix}$$

- R in 2 dimensional space (keep both components):

$$R = \Sigma V^T = \begin{bmatrix} 13.93 & 0 \\ 0 & 0.92 \end{bmatrix} \begin{bmatrix} -0.62 & -0.79 \\ 0.79 & -0.62 \end{bmatrix} = \begin{bmatrix} -8.63 & -10.94 \\ 0.72 & -0.57 \end{bmatrix}$$

$$R_0 = \begin{bmatrix} -8.63 \\ 0.72 \end{bmatrix} \quad R_1 = \begin{bmatrix} -10.94 \\ -0.57 \end{bmatrix} \quad \text{dist}(R_0, R_1) = 2.646$$

Principal Component Analysis DEMO

Jupyter Notebook for demo:

- UnsupervisedML/Examples/Section09/PCA.ipynb

Course Resources at:

- <https://github.com/satishchandrareddy/UnsupervisedML/>

Unsupervised Machine Learning with Python

Section 9.2: Principal Component Analysis Code Design

PCA Code Design

- This section contains design information for the PCA Code based on algorithm presented in Section 9.1
- Code located at `UnsupervisedML/Code/Programs`
- Stop video here, if you would like to do code design yourself

PCA Code Design

Create PCA class with methods for:

- Performing SVD
- Computing dataset in reduced number of dimensions
- Computing reconstructed dataset
- Plotting cumulative variance proportion

pca class: Principal Variables

Variable	Type	Description
self.Xmean	2d numpy array	X_{mean} of dataset (column vector of dimension $d \times 1$)
self.dimension	integer	Number of dimensions in dataset
self.nsample	Integer	Number of data points
self.U	2d numpy array	U from SVD of $X - X_{mean}$ ($d \times N$) $N = \min(d, M)$
self.Sigma	1d numpy array	Singular values of $X - X_{mean}$ (N entries)
self.Vt	2d numpy array	V^T from SVD of $X - X_{mean}$ ($N \times M$)
self.cumulative_ variance_ proportion	1d numpy array	Cumulative variance proportion (N entries)

pca class – Key Methods

Method	Input	Description
<code>__init__</code>		Constructor for pca class
<code>fit</code>	X (2d np.array)	Computes compact SVD of X – Xmean and stores results in self.U, self.Sigma, and self.Vt Return: nothing
<code>get_dimension</code>	variance_capture (float)	Computes minimum number of principal components to retain at least variance_capture proportion of total variance Return: number of principal components K
<code>data_reduced_dimension</code>	**kwargs reduced_dim (integer) variance_capture (float)	Computes the dataset in U coordinate system with reduced number of dimensions. User specifies either reduced_dim directly or the proportion of variance to be captured. Return: R (2d numpy array)
<code>data_reconstructed</code>	**kwargs reduced_dim (integer) variance_capture (float)	Computes the reconstructed dataset in original coordinate system with reduced number of dimensions. User specifies either reduced_dim directly or the proportion of variance to be captured. Return: X_R (2d numpy array)
<code>plot_cumulative_variance_proportion</code>		Plots the cumulative variance proportion Return: nothing

Unsupervised Machine Learning with Python

Section 9.3: Principal Component Analysis Code Walkthrough

PCA Code Walkthrough

Code located at:

- UnsupervisedML/Code/Programs

Files to Review	Description
pca.py	class for principal component analysis

Course Resources at:

- <https://github.com/satishchandrareddy/UnsupervisedML/>
- Stop video if you would like to implement code yourself first

Unsupervised Machine Learning with Python

Section 9.4: PCA Applied to MNIST Digits Dataset

MNIST Digits Dataset

- Thousands of handwritten digit images with 28x28 resolution
- Data Source: <http://yann.lecun.com/exdb/mnist/>
- Used in machine learning research and teaching



Collage of 160 individual digit images

By Josef Steppan - Own work, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=64810040>

MNIST Digits – Format of Data Files

- Each row contains data for one image
 - First column is the digit label (0,1,...,9)
 - Columns 2 – 785 are the intensities (integers between 0=white and 255=black)
 - 784 dimensions
 - Take transpose to convert feature matrix to (dimension x sample) format
 - Standard practice is to divide pixel values by 255 so between 0 and 1

Normal

Page Break Preview

Page Layout

Custom Views

☒ Ruler

☒ Formula Bar

☒ Gridlines

☒ Headings

Zoom

100%

Zoom to Selection

New Window

Arrange All

Freeze Panes

Split

Hide

Unhide

View Side by Side

Synchronous Scrolling

Reset Window Position

Switch Windows

Macros

Workbook Views

Show

Zoom

Window

Macros

POSSIBLE DATA LOSS

Some features might be lost if you save this workbook in the comma-delimited (.csv) format. To preserve these features, save it in an Excel file format.

Don't show again

Save As...

EJ1

✕

✓

fx


pixel138

	A	DT	DU	DV	DW	DX	DY	DZ	EA	EB	EC	ED	EE	EF	EG	EH	EI	EJ	EK	EL	p
1	label	pixel122	pixel123	pixel124	pixel125	pixel126	pixel127	pixel128	pixel129	pixel130	pixel131	pixel132	pixel133	pixel134	pixel135	pixel136	pixel137	pixel138	pixel139	pixel140	p
2	7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
3	2	253	253	253	253	253	253	218	30	0	0	0	0	0	0	0	0	0	0	0	
4	1	0	0	0	0	0	0	38	254	109	0	0	0	0	0	0	0	0	0	0	
5	0	0	0	11	150	253	202	31	0	0	0	0	0	0	0	0	0	0	0	0	
6	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
7	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
8	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
9	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
10	5	0	0	0	0	0	0	0	17	47	47	47	16	129	85	47	0	0	0	0	
11	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
12	0	0	61	3	42	118	193	118	118	61	0	0	0	0	0	0	0	0	0	0	
13	6	150	252	252	125	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
14	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	

Example MNIST Digits

Original Dataset: 60000 images
Plot 25 randomly chosen images

Images of Sample MNIST Digits




PCA for MNIST Digits Dataset

Perform PCA for MNIST Digits Dataset

- Dataset X: $d = 784$ dimensions and $M = 60000$ images
- Compute compact SVD of $\bar{X} = X - X_{mean} = U\Sigma V^T$

$$\bar{X} (d \times M) = U (d \times d) \Sigma (d \times d) V^T (d \times M)$$

Cumulative Variance Proportion




PCA for MNIST: Original and Reconstructed

Original Dataset:

Plot 25 randomly chosen images


Images of Sample MNIST Digits



Reconstructed Dataset:

90% variance capture (87 dimensions)

Images of Sample MNIST Digits



mnist class Code Design

Method	Input	Description
__init__		Constructor for mnist class – saves data directory Return: nothing
load_train	nsample (integer)	Loads nsample (maximum 60,000) images and corresponding class labels from the train dataset Return: X (2d numpy array), class_label (1d numpy array)
load_valid	nsample (integer)	Loads nsample (maximum 10,000) images and corresponding class labels from the validation dataset Return: X (2d numpy array), class_label (1d numpy array)
plot	X (2d numpy array) seed (integer)	Creates figure of 25=5x5 images randomly chosen from the dataset X. seed is used to set up the seed for the random number generator. Return: nothing

PCA for MNIST Digits Code Walkthrough

Programs and data located at

- UnsupervisedML/Code/Programs
- UnsupervisedML/Code/Data_MNIST

Files to Review	Description
Programs/data_mnist.py	Class for loading and plotting mnist digits dataset
Programs/driver_pca_mnist.py	Driver for performing pca for mnist dataset
Data_MNIST/MNIST_valid_10K.csv	Validation dataset (10000 images)

Course Resources at:

- <https://github.com/satishchandrareddy/UnsupervisedML/>
- Stop video if you would like to implement code yourself first

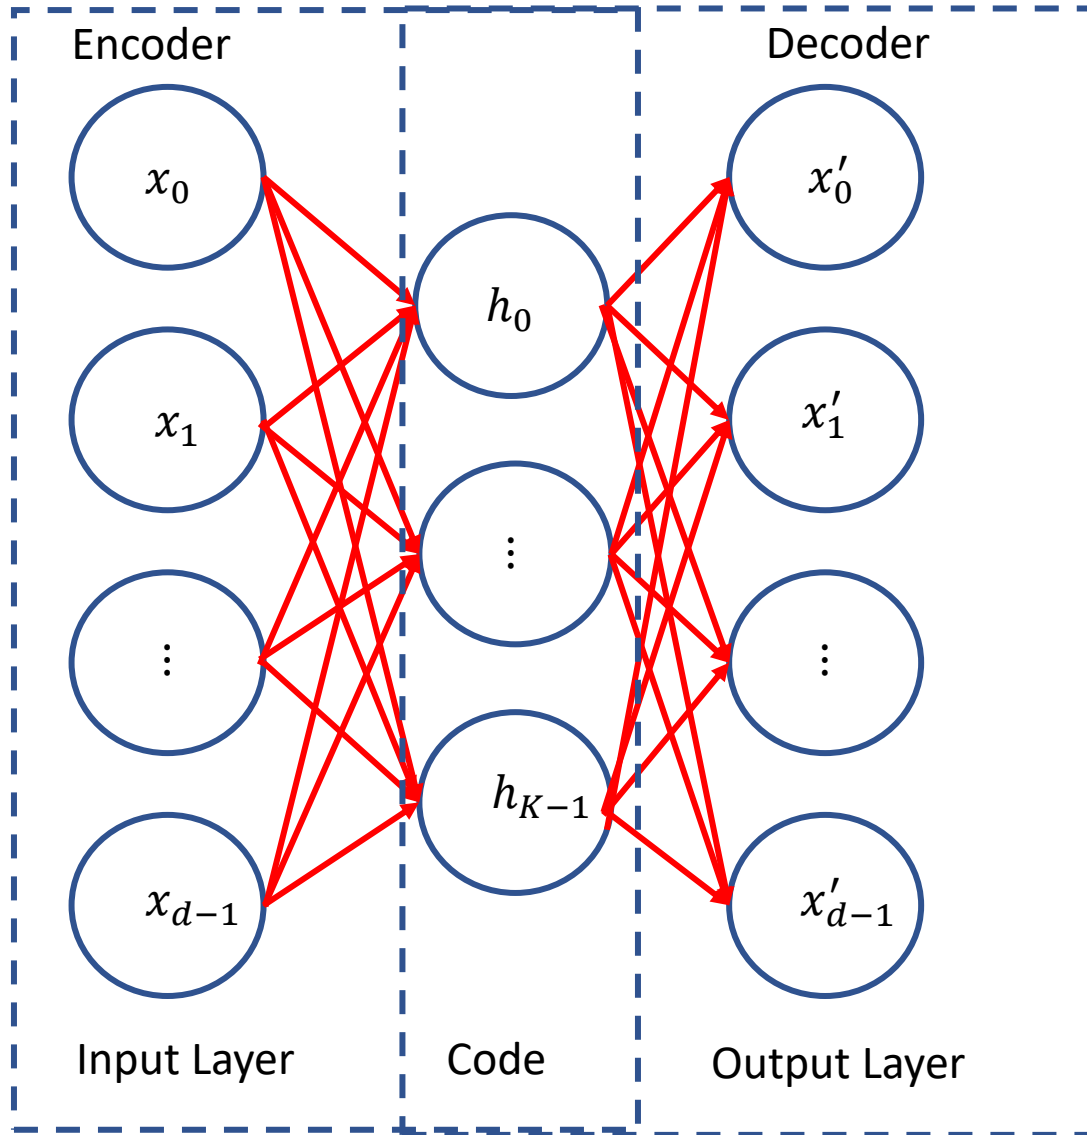
Unsupervised Machine Learning with Python

Section 9.5: Autoencoders

What is an Autoencoder?

- An Autoencoder is type of Artificial Neural Network for learning efficient representations of the feature vector
- Can be used to reduce dimensions
- Whereas PCA (based on SVD) is linear, Autoencoder can use a nonlinear approach

Sample Autoencoder Neural Network



- Input: (d dimensions)

$$X = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{d-1} \end{bmatrix}$$

- Encoder: Map X to H (K dimensions). “Code” H is a new feature vector representation of X

$$H = \begin{bmatrix} h_0 \\ \vdots \\ h_{K-1} \end{bmatrix}$$

- Decoder: Map H to X' (d dimensions). X' is a reconstructed version of X

$$X' = \begin{bmatrix} x'_0 \\ x'_1 \\ \vdots \\ x'_{d-1} \end{bmatrix}$$

Autoencoder: Determining Mappings

- Encoder:

$$H = f(WX + b)$$

X (d dimensions) is data point, W is Kxd matrix, b is Kx1 vector, f is activation function, H is (K dimensions) reduced dimension version of data point

- Decoder:

$$X' = g(W'H + b')$$

X' is (d dimensions) reconstructed data point, W' is dxK matrix, b' is dx1 vector, g is activation function

- Given a dataset of M datapoints X_0, X_1, \dots, X_{M-1} and pre-specified activation functions f, g: Determine unknown W, b, W', b' by minimizing mean squared error loss function:

$$Loss = \frac{1}{M} \sum_{i=0}^{M-1} dist(X'_i, X_i)^2$$

Autoencoder: Notes

- Autoencoder set up is similar to Supervised Learning
 - Key difference is that there is no label Y for each data point
- Determine W, b, W', b' using optimizer such as Gradient Descent and backpropagation to determine derivatives
- Can use familiar activation functions such as sigmoid, Relu for f and g
- “Code” representation H is related nonlinearly to X (if nonlinear activation functions used)
- Code representation H is equivalent to reduced dimension version R from Principal Component Analysis