

Unsupervised Machine Learning

Section 9.1: Dimension Reduction

Dimension Reduction

- Machine learning problems may involve datasets in 100s or 1000s of dimensions
- More dimensions generally means slower computation
- Dimension Reduction attempts to map feature vectors into a lower dimensional space while retaining as much information as possible

Dimension Reduction

Two Approaches in this Section:

- Principal Component Analysis (PCA):
 - “Linear” approach
 - Cool application of singular value decomposition
- Autoencoding:
 - “Nonlinear” approach
 - Uses techniques from supervised learning to learn new representation of data

Section 9.2: Principal Component Analysis: Algorithm


PCA and SVD

- SVD is used to find a new basis u_0, u_1, \dots, u_{d-1}
- Relevance of each basis vector (or component) is determined by its corresponding singular value, which are ordered in decreasing value
- Idea is to project data onto space spanned by first K basis vectors u_0, u_1, \dots, u_{K-1} , choosing K to retain sufficient information

Example: 4 Data Points in 2D

- Consider 4 data points in 2 dimensions

$$X = [X_0 \quad X_1 \quad X_2 \quad X_3] = \begin{bmatrix} 1 & 1.6 & -0.5 & -1.6 \\ 0.9 & 1.65 & -0.6 & -1.5 \end{bmatrix}$$



SVD of Data Set

- Compute SVD of X



$$X = U\Sigma V^T$$

$$X = U\Sigma V^T = [u_0 \quad u_1] \begin{bmatrix} \sigma_0 & 0 & 0 & 0 \\ 0 & \sigma_1 & 0 & 0 \end{bmatrix} \begin{bmatrix} v_0^T \\ v_1^T \\ v_2^T \\ v_3^T \end{bmatrix}$$

$$X = \begin{bmatrix} 0.71 & -0.70 \\ 0.70 & -0.71 \end{bmatrix} \begin{bmatrix} 3.54 & 0 & 0 & 0 \\ 0 & 0.12 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0.38 & 0.65 & -0.22 & -0.62 \\ -0.47 & 0.46 & -0.63 & 0.41 \\ -0.26 & 0.58 & 0.74 & 0.19 \\ 0.75 & 0.16 & -0.03 & 0.64 \end{bmatrix}$$

New Coordinate system with basis u_0 and u_1

- Consider new coordinate system using u_0 and u_1 as basis



Data points in u0-u1 Coordinate System

- Original X

$$X = \begin{bmatrix} 1 & 1.6 & -0.5 & -1.6 \\ 0.9 & 1.65 & -0.6 & -1.5 \end{bmatrix}$$

- Data points in u0 and u1 coordinate system

$$X_U = \Sigma V^T = \begin{bmatrix} 3.54 & 0 & 0 & 0 \\ 0 & 0.12 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0.38 & 0.65 & -0.22 & -0.62 \\ -0.47 & 0.46 & -0.63 & 0.41 \\ -0.26 & 0.58 & 0.74 & 0.19 \\ 0.75 & 0.16 & -0.03 & 0.64 \end{bmatrix}$$

$$X_U = \begin{bmatrix} 1.34 & 2.30 & -0.78 & -2.20 \\ -0.06 & -0.06 & -0.08 & 0.05 \end{bmatrix}$$

← Row0 gives units of u0 for each data point

← Row1 gives units of u1 for each data point

Reduce Dimensions

- Only keep information in u0 direction (1 dimension)
- Data points in reduced number of dimensions

$$R = [\sigma_0][v_0^T] = [3.54][0.38 \quad 0.65 \quad -0.22 \quad -0.62] = [1.34 \quad 2.30 \quad -0.78 \quad -2.20]$$

Units of u0 for each data point



Reconstruct in Original Space


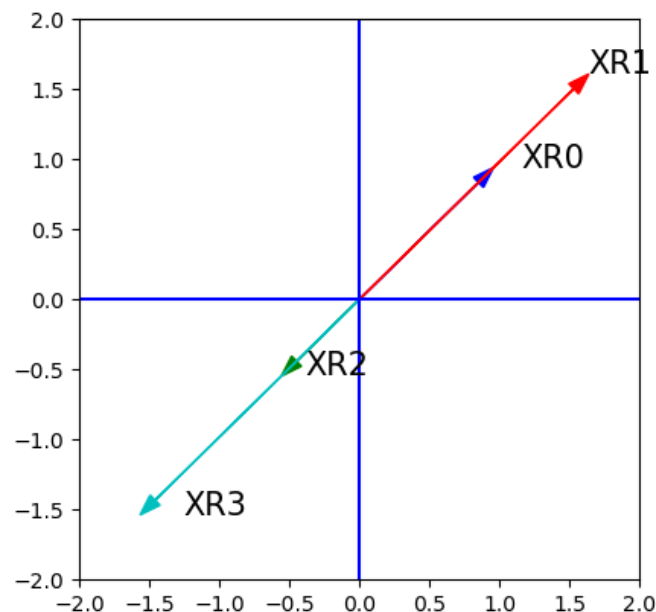
- Can reconstruct data points in 2d space using information in u_0 direction only

$$X_R = [u_0][\sigma_0][v_0^T]$$

$$X_R = \begin{bmatrix} 0.71 \\ 0.70 \end{bmatrix} [3.54] \begin{bmatrix} 0.38 & 0.65 & -0.22 & -0.62 \end{bmatrix}$$

$$X_R = \begin{bmatrix} 0.96 & 1.64 & -0.55 & -1.56 \\ 0.94 & 1.61 & -0.54 & -1.54 \end{bmatrix}$$

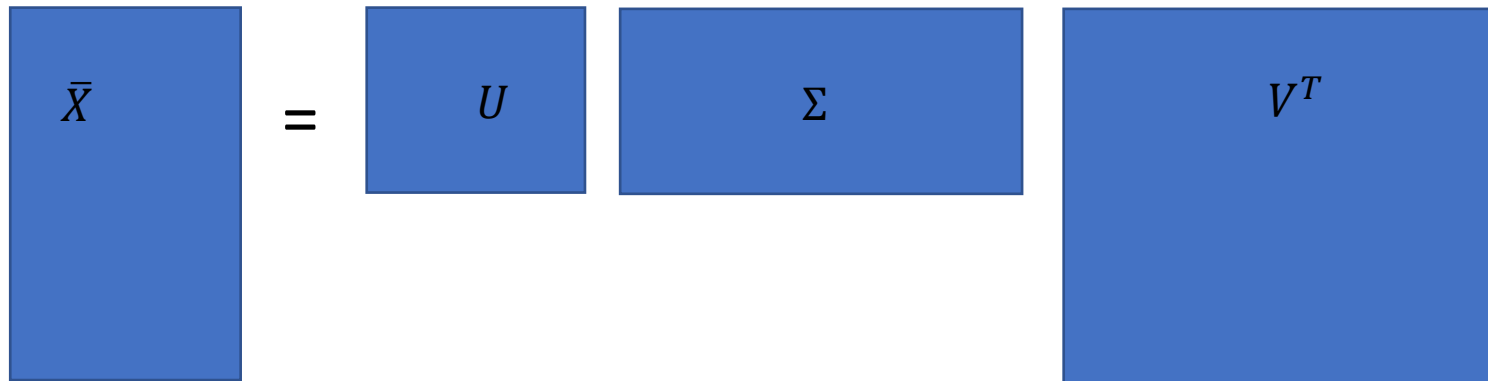
$$X = \begin{bmatrix} 1 & 1.6 & -0.5 & -1.6 \\ 0.9 & 1.65 & -0.6 & -1.5 \end{bmatrix}$$



PCA Algorithm

PCA Algorithm:

- Start with data matrix X (number of features x number of samples)
- (1) Subtract mean of data $\bar{X} = X - X_{mean}$ (translate data points by same amount)
 - (2) Compute the singular value decomposition of $\bar{X} = X - X_{mean} = U\Sigma V^T$




A diagram illustrating the Singular Value Decomposition (SVD) equation. It consists of four blue rectangular boxes arranged horizontally, separated by an equals sign. The first box on the left is tall and contains the symbol \bar{X} . The second box is short and contains the symbol U . The third box is short and contains the symbol Σ . The fourth box is tall and contains the symbol V^T .

$$\bar{X} = U \Sigma V^T$$

PCA Algorithm

(3) Pick K principal components



(first K columns of U , first K diagonal entries of Σ , first K rows of V^T)

(4) Data points in K dimensional (u_0, \dots, u_{K-1}) coordinate system given by:

$$R = \begin{bmatrix} \sigma_0 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_{K-1} \end{bmatrix} \begin{bmatrix} v_0^T \\ \vdots \\ v_{K-1}^T \end{bmatrix}$$

PCA Algorithm

(5) Reconstructed \bar{X} (keeping only K principal components) in original space

$$\bar{X}_R = [u_0 \quad \cdots \quad u_{K-1}] \begin{bmatrix} \sigma_0 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_{K-1} \end{bmatrix} \begin{bmatrix} v_0^T \\ \vdots \\ v_{K-1}^T \end{bmatrix}$$

(6) Add back mean to get reconstructed X (keeping only K principal components)

$$X_R = \bar{X}_R + X_{mean}$$

Principal Component Analysis: Choosing K

- Covariance matrix for data set:

$$Cov = \frac{1}{M} (X - X_{mean})(X - X_{mean})^T$$

- Total variance is sum of eigenvalues of Covariance matrix, which is same as sum of square of singular values of $\frac{1}{\sqrt{M}} (X - X_{mean})$
- Define


$$Variance(K) = \frac{1}{M} \sum_{k=0}^{K-1} \sigma_k^2$$

- Instead of specifying K directly, choose smallest number of principal components so that $Variance(K)/Variance(M)$ is at least as large as specified percentage

Section 9.3: PCA for MNIST Dataset

MNIST Digits Dataset

- Thousands of handwritten digit images with 28x28 resolution
- Data Source: <http://yann.lecun.com/exdb/mnist/>
- Used extensively for testing machine learning algorithms



Collage of 160 individual digit images

By Josef Steppan - Own work, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=64810040>


MNIST Digits – Format of Data Files

- Each row represents label and intensities for one image
 - First column is the digit label (0,1,...,9)
 - Columns 2 – 785 are the intensities
 - Take transpose to convert feature matrix and to correct format
 - Standard practice is to divide pixel values by 255 so between 0 and 1


	A	DT	DU	DV	DW	DX	DY	DZ	EA	EB	EC	ED	EE	EF	EG	EH	EI	EJ	EK	EL	p
1	label	pixel122	pixel123	pixel124	pixel125	pixel126	pixel127	pixel128	pixel129	pixel130	pixel131	pixel132	pixel133	pixel134	pixel135	pixel136	pixel137	pixel138	pixel139	pixel140	p
2	7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
3	2	253	253	253	253	253	253	218	30	0	0	0	0	0	0	0	0	0	0	0	0
4	1	0	0	0	0	0	0	38	254	109	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	11	150	253	202	31	0	0	0	0	0	0	0	0	0	0	0	0	0
6	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
7	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
8	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
9	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
10	5	0	0	0	0	0	0	0	17	47	47	47	16	129	85	47	0	0	0	0	0
11	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
12	0	0	61	3	42	118	193	118	118	61	0	0	0	0	0	0	0	0	0	0	0
13	6	150	252	252	125	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
14	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

PCA for MNIST: Original and Reconstructed


Original
Dataset: 6000 image
784 dimensions
Plot of 25 images




Reconstructed
99.9% Variance
476 dimensions



Reconstructed
99% Variance
323 dimensions



Reconstructed
90% Variance
84 dimensions



Section 9.4: PCA: Code Design

PCA Code Design

- This section contains design information for the PCA Code
- Design based on theory in Section 10.1
- Code located at `UnsupervisedML/Code/Clustering/pca.py`
- Stop video here, if you would like to do code design yourself

PCA Code Design: To Do

Component	Description
class pca	class for Principal Component Analysis
driver_pca	driver for pca example
load_mnist	Function for loading mnist_data set

pca class: Principal Variables

Variable	Type	Description
self.U	2d numpy array	U from SVD of $X - X_{mean}$
self.Sigma	1d numpy array	Singular values of $X - X_{mean}$
self.Vh	2d numpy array	V^T from SVD of $X - X_{mean}$

pca class – Key Methods

Method	Input	Description
<code>__init__</code>		Constructor for pca class
<code>fit</code>	X (2d np.array)	Computes SVD of X – Xmean and stores results in self.U, self.Sigma, and self.vh Return: nothing
<code>get_dimension</code>	variance_capture (float)	Computes number of principal components to capture specified proportion of variance Return: number of principal components K
<code>data_reduced_dimension</code>	**kwargs reduced_dim (integer) variance_capture (float)	Computes the dataset in the reduced number of dimensions. User specifies either reduced_dim directly or the proportion of variance to be captured. Return: R (2d numpy array)
<code>data_reconstructed</code>	**kwargs reduced_dim (integer) variance_capture (float)	Computes the reconstructed dataset using only K principal components. User specifies either reduce_dim directly or the proportion of variance to be captured. Return: X_R

Section 9.5: PCA: Code Walkthrough

PCA: Code Walkthrough

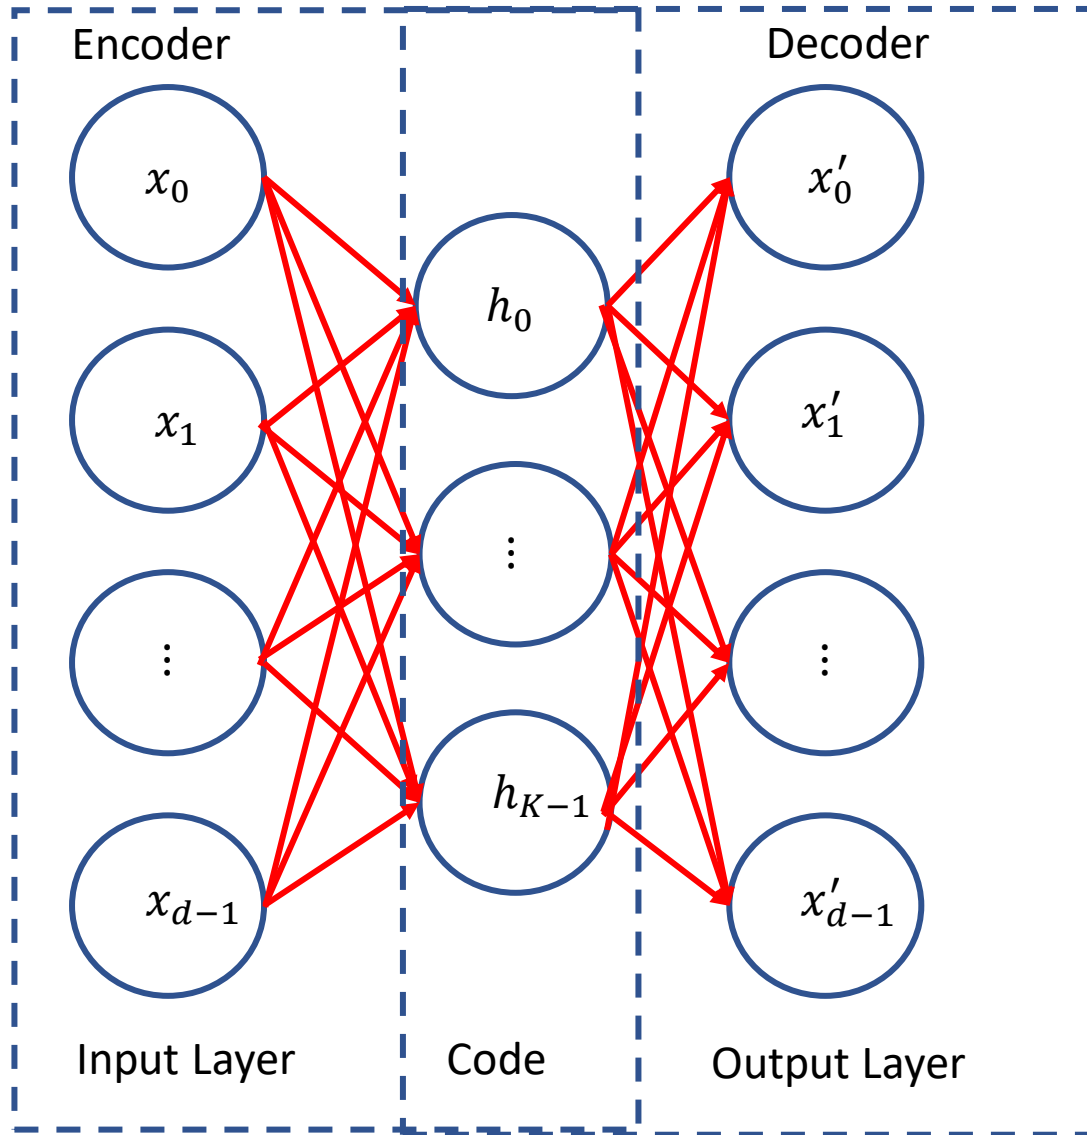
- Code is located in:
Folder: UnsupervisedML/Code/Clustering
Files: pca.py, driver_pca.py, load_mnist.py
- Stop video here, if you would like to do coding yourself before seeing my implementation

Section 9.6: Autoencoders

What is an Autoencoder?

- An Autoencoder is type of Artificial Neural Network for learning efficient representations of the feature vector in unsupervised manner
- Can be used to reduce dimensions
- Whereas PCA is based on SVD is linear, Autoencoder is a nonlinear approach

Sample Autoencoder Neural Network



- Input: (d dimensions)

$$X = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{d-1} \end{bmatrix}$$

- Encoder: Map X to H (K dimensions). “Code” H is a new feature vector representation of X

$$H = \begin{bmatrix} h_0 \\ \vdots \\ h_{K-1} \end{bmatrix}$$

- Decoder: Map H to X' (d dimensions). X' is a reconstructed version of X

$$X' = \begin{bmatrix} x'_0 \\ x'_1 \\ \vdots \\ x'_{d-1} \end{bmatrix}$$

Autoencoder: Determining Mappings

- Encoder:

$$H = f(WX + b)$$

W is $K \times d$ matrix, b is $K \times 1$ vector, f is activation function

- Decoder:

$$X' = g(W'H + b')$$

W' is $d \times K$ matrix, b' is $d \times 1$ vector, g is activation function

- Given a dataset of $M-1$ datapoints: X_0, X_1, \dots, X_{M-1} for fixed activation functions f , and g , determine unknown W, b, W', b' by minimizing mean squared error loss function:

$$Loss = \frac{1}{M} \sum_{i=0}^{M-1} dist(X'_i, X_i)^2$$

Autoencoder: Notes

- Autoencoder set up is similar to Supervised Learning
 - Key difference is that there is no label Y for each data point
- Determine W, b, W', b' using optimizer such as Gradient Descent and backpropagation to determine derivatives
- Can use familiar activation functions such as sigmoid, Relu for f and g
- “Code” representation H is related nonlinearly to X (if nonlinear activation functions used)
- Use code representation H in unsupervised learning algorithm instead of original X